

One claim from the prompt

“One ball is tossed in, triggering a cascade of 1,000 traps.”

Two separable things are asserted: that the cascade happens at all, and how many traps it involves.

```
check(c1, judge("the cascade occurs") and eq(count("mousetrap"), 1000))
```

Written before a single frame is read.
The numeral is what adds the second test.

one claim, two measurements

Does the cascade occur?

Goes to the general verifier, which reads the frames densely. Occurrence questions go there because a specialist asked for something absent may still hand back its best candidate.

Returns an occurrence decision.

evidence only

How many traps are there?

Goes to the counting specialist, which returns the number of traps it can see and nothing else — no comparison against the requested figure, and no verdict of its own.

Returns a count.

evidence only

How the two outcomes compose

any part contradicted → the claim is contradicted

every part supported → supported | anything else → unresolved

A count far below the requested number contradicts the claim on its own, even when the cascade plainly occurs. A claim carrying only an occurrence check could not say that.

Why this claim is in the benchmark

An annotator marked exactly this prompt fragment, wrote that the video shows 63 mousetraps rather than the thousand requested, filed it under counting, and rated it 3 out of 5 for severity — the only flaw on that clip. What the critic returns on this clip is not reported here.

Gray fill marks a learned component; a white box with an outline is a deterministic rule; a dashed box is text we were handed rather than anything the critic produced.