


The same depth control in all 18 clips; only the prompt’s motion clause changes



● one mark = one clip    ----- what the control video itself does (0.87)

Dashed line: the control clip measured with the same metric — its object crosses 0.87 of the frame width before leaving frame.

This is the motion the clips were conditioned on. Under each panel: the largest change in travel from changing only the wording, and from changing only the seed, each averaged over that panel’s other factor.

“straight down” is the adversarial wording: it asks for the one thing the control video does not do